

Package ‘Rgtsvm’

March 19, 2017

Version 0.31

Date 2016-04-26

Title The SVM package for GPU architecture based on the GTSVM software and e1071 package.

Imports graphics, grDevices, class, bit64, tools, methods

Suggests SparseM, Matrix

Description Training and prediction of SVM based on GTSVM are available on GPU architecture.

License GPLv3

LazyLoad yes

Author Zhong Wang <zw355@cornell.edu>

Maintainer Zhong Wang <zw355@cornell.edu>

NeedsCompilation yes

Repository CRAN

R topics documented:

load.svmLight	1
plot.gtsvm	2
plot.tune	3
predict.gtsvm	4
svm	5
tune	9
tune.control	11

Index	13
--------------	-----------

load.svmLight	<i>Load SVMlight data file into a sparse matrix.</i>
---------------	--

Description

Load SVMlight data file into a sparse matrix.

Usage

```
load.svmLight(filename)
```

Arguments

filename SVM light filename.

Details

The file must be svmlight format.(<http://svmlight.joachims.org/>)

Value

A sparse matrix is returned if the file is loaded or downloaded successfully.

Author(s)

Zhong Wang (R interface) <zw355@cornell.edu>

Examples

```
mat <-load.svmlight("http://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/multiclass/glass.scale")
str(mat);
```

plot.gtsvm

Plot SVM Objects

Description

Generates a scatter plot of the input data of a svm fit for classification models by highlighting the classes and support vectors. Optionally, draws a filled contour plot of the class regions.

Usage

```
## S3 method for class 'svm'
plot(x, data, formula, fill = TRUE, grid = 50, slice = list(),
     symbolPalette = palette(), svSymbol = "x", dataSymbol = "o", ...)
```

Arguments

x	An object of class svm
data	data to visualize. Should be the same used for fitting.
formula	formula selecting the visualized two dimensions. Only needed if more than two input variables are used.
fill	switch indicating whether a contour plot for the class regions should be added.
grid	granularity for the contour plot.
slice	a list of named values for the dimensions held constant (only needed if more than two variables are used). The defaults for unspecified dimensions are 0 (for numeric variables) and the first level (for factors). Factor levels can either be specified as factors or character vectors of length 1.
symbolPalette	Color palette used for the class the data points and support vectors belong to.
svSymbol	Symbol used for support vectors.
dataSymbol	Symbol used for data points (other than support vectors).
...	additional graphics parameters passed to filled.contour and plot.

Author(s)

David Meyer
 <David.Meyer@R-project.org>

See Also

[svm](#)

Examples

```
## a simple example
data(cats, package = "MASS")
m <- svm(Sex~., data = cats)
plot(m, cats)

## more than two variables: fix 2 dimensions
data(iris)
m2 <- svm(Species~., data = iris)
plot(m2, iris, Petal.Width ~ Petal.Length,
      slice = list(Sepal.Width = 3, Sepal.Length = 4))

## plot with custom symbols and colors
plot(m, cats, svSymbol = 1, dataSymbol = 2, symbolPalette = rainbow(4),
      color.palette = terrain.colors)
```

plot.tune

Plot Tuning Object

Description

Visualizes the results of parameter tuning.

Usage

```
## S3 method for class 'tune'
plot(x, type = c("contour", "perspective"), theta = 60,
      col = "lightblue", main = NULL, xlab = NULL, ylab = NULL,
      swapxy = FALSE, transform.x = NULL, transform.y = NULL,
      transform.z = NULL, color.palette = hsv_palette(),
      nlevels = 20, ...)
```

Arguments

x	an object of class tune
type	choose whether a contour plot or a perspective plot is used if two parameters are to be visualized. Ignored if only one parameter has been tuned.
theta	angle of azimuthal direction.
col	the color(s) of the surface facets. Transparent colors are ignored.

main	main title
xlab, ylab	titles for the axes. N.B. These must be character strings; expressions are not accepted. Numbers will be coerced to character strings.
swapxy	if TRUE, the parameter axes are swaped (only used in case of two parameters).
transform.x, transform.y, transform.z	functions to transform the parameters (x and y) and the error measures (z). Ignored if NULL.
color.palette	color palette used in contour plot.
nlevels	number of levels used in contour plot.
...	Further graphics parameters.

Author(s)

David Meyer (based on C/C++-code by Chih-Chung Chang and Chih-Jen Lin)
 <David.Meyer@R-project.org>

See Also

[tune](#)

Examples

```
data(iris)
obj <- tune.svm(Species~., data = iris, sampling = "fix",
               gamma = 2^c(-8,-4,0,4), cost = 2^c(-8,-4,-2,0))
plot(obj, transform.x = log2, transform.y = log2)
plot(obj, type = "perspective", theta = 120, phi = 45)
```

predict.gtsvm

Predict Method for Support Vector Machines

Description

This function predicts values based upon a model trained by `svm` in package *Rgtsvm*.

Usage

```
## S3 method for class 'gtsvm'
predict(object, newdata,
        decision.values = FALSE,
        probability = FALSE,
        verbose=FALSE,
        ...,
        na.action = na.omit)
```

Arguments

<code>object</code>	Object of class "gtsvm", created by <code>svm</code> in <i>Rgtsvm</i> package.
<code>newdata</code>	An object containing the new input data: either a matrix or a sparse matrix (object of class <code>Matrix</code> provided by the Matrix package, or of class <code>matrix.csr</code> provided by the SparseM package, or of class <code>simple_triplet_matrix</code> provided by the slam package). A vector will be transformed to a $n \times 1$ matrix.
<code>decision.values</code>	Logical controlling whether the decision values of all binary classifiers computed in multiclass classification shall be computed and returned.
<code>probability</code>	Logical indicating whether class probabilities should be computed and returned. Only possible if the model was fitted with the <code>probability</code> option enabled.
<code>verbose</code>	logical value indicating whether some algorithm information (default: FALSE)
<code>...</code>	Currently not used.
<code>na.action</code>	A function to specify the action to be taken if 'NA's are found. The default action is <code>na.omit</code> , which leads to rejection of cases with missing values on any required variable. An alternative is <code>na.fail</code> , which causes an error if NA cases are found. (NOTE: If given, this argument must be named.)

Value

A vector of predicted values (for classification: a vector of labels, for density estimation: a logical vector). If `score` is TRUE, the vector gets a "decision.values" attribute containing a $n \times c$ matrix (n number of predicted values, c number of classifiers) of all c binary classifiers' decision values.

Note

If the training set was scaled by `svm` in *Rgtsvm*, the new data is scaled accordingly using scale and center of the training data.

Author(s)

Zhong Wang (R interface & epe-regression in CUDA) <zw355@cornell.edu>

David Meyer (R interface in e1071) <David.Meyer@R-project.org>

Andrew Cotter, Nathan Srebro ,Joseph Keshet (C/C++ code in CUDA) <http://ttic.uchicago.edu/textasciitildecotter/project>

See Also

[svm](#)

Description

`svm` in *Rgtsvm* package is used to train a support vector machine by the C-classification and epsilon regression method. A formula interface is provided.

Usage

```
## S3 method for class 'formula'
svm(formula, data = NULL, ..., na.action = na.omit, scale = TRUE);

## Default S3 method:
svm(x,
     y          = NULL,
     scale      = TRUE,
     type       = "C-classification",
     kernel     = "radial",
     degree     = 3,
     gamma      = 0.05,
     coef0      = 0,
     cost       = 1,
     class.weights = NULL,
     tolerance  = 0.001,
     epsilon    = 0.1,
     shrinking  = TRUE,
     cross      = 0,
     probability = FALSE,
     fitted     = TRUE,
     rough.cross = 0,
     no.change.x = TRUE,
     verbose    = FALSE,
     ...,
     subset,
     na.action = na.omit)
```

Arguments

formula	a symbolic description of the model to be fit.
data	an optional data frame containing the variables in the model. By default the variables are taken from the environment which ‘svm’ is called from.
x	a data matrix, a vector, or a sparse matrix (object of class Matrix provided by the Matrix package, or of class matrix.csr provided by the SparseM package, or of class simple_triplet_matrix provided by the slam package).
y	a response vector with one label for each row/component of x. Can be either a factor (for classification tasks) or a numeric vector (for regression).
scale	A logical vector indicating the variables to be scaled. If scale is of length 1, the value is recycled as many times as needed. Per default, data are scaled internally (both x and y variables) to zero mean and unit variance. The center and scale values are returned and used for later predictions.
type	only C-classification or eps-regression available.
kernel	the kernel used in training and predicting. You might consider changing some of the following parameters, depending on the kernel type.

linear: $u'v$

polynomial: $(\gamma u'v + \text{coef0})^{\text{degree}}$

radial basis: $e^{(\gamma |u - v|^2)}$

sigmoid: $\tanh(\gamma u'v + \text{coef0})$

degree	parameter needed for kernel of type polynomial (default: 3)
gamma	parameter needed for all kernels except linear (default: 1/(data dimension))
coef0	parameter needed for kernels of type polynomial and sigmoid (default: 0)
cost	cost of constraints violation (default: 1)—it is the ‘C’-constant of the regularization term in the Lagrange formulation.
class.weights	a named vector of weights for the different classes, used for asymmetric class sizes. Not all factor levels have to be supplied (default weight: 1). All components have to be named.
tolerance	tolerance of termination criterion (default: 0.001)
epsilon	epsilon in the insensitive-loss function (default: 0.1)
shrinking	option whether to use the shrinking-heuristics (default: TRUE)
cross	if a integer value $k > 0$ is specified, a k -fold cross validation on the training data is performed to assess the quality of the model: the accuracy rate for classification and the Mean Squared Error for regression
fitted	logical indicating whether the fitted values should be computed and included in the model or not (default: TRUE)
probability	logical indicating whether the model should allow for probability predictions.
rough.cross	Number which is less than cross, indicating how many tests are performed for cross-validation. The function will return partial tests for cross-validation rather than all repeated tests in order to reduce the running time.
no.change.x	Logical value indicating whether the function can change the x parameter. If x is a big matrix, it would be save memory to use this parameter FALSE.
verbose	logical value indicating whether some algorithm information (default: FALSE)
...	additional parameters for the low level fitting function <code>svm.default</code>
subset	An index vector specifying the cases to be used in the training sample. (NOTE: If given, this argument must be named.)
na.action	A function to specify the action to be taken if NAs are found. The default action is <code>na.omit</code> , which leads to rejection of cases with missing values on any required variable. An alternative is <code>na.fail</code> , which causes an error if NA cases are found. (NOTE: If given, this argument must be named.)

Details

Rgtsvm internally uses a sparse matrix and regular matrix.

If the predictor variables include factors, the formula interface must be used to get a correct model matrix.

`plot.gtsvm` allows a simple graphical visualization of classification models.

Value

An object of class `"gtsvm"` containing the fitted model, including:

SV	The resulting support vectors (possibly scaled).
index	The index of the resulting support vectors in the data matrix. Note that this index refers to the preprocessed data (after the possible effect of <code>na.omit</code> and <code>subset</code>)
coefs	The corresponding coefficients times the training labels.

Author(s)

Zhong Wang (R interface & epe-regression in CUDA) <zw355@cornell.edu>

David Meyer (R interface in e1071) <David.Meyer@R-project.org>

Andrew Cotter, Nathan Srebro, Joseph Keshet (C/C++ code in CUDA) <http://ttic.uchicago.edu/textasciitildecotter/proje>

References

- Andrew Cotter, Nathan Srebro, Joseph Keshet. "A GPU-Tailored Approach for Training Kernelized SVMs". 17th ACM SIGKDD Conference on Knowledge Discovery and Data Mining. 2011.
- Chang, Chih-Chung and Lin, Chih-Jen:
LIBSVM: a library for Support Vector Machines
<http://www.csie.ntu.edu.tw/~cjlin/libsvm>

See Also

[predict.gtsvm](#) [plot.gtsvm](#) [matrix.csr](#) (in package **SparseM**)

Examples

```
data(iris)
attach(iris)

## classification mode
# default with factor response:
model <- svm(Species ~ ., data = iris)

# alternatively the traditional interface:
x <- subset(iris, select = -Species)
y <- Species
model <- svm(x, y)

print(model)
summary(model)

# test with train data
pred <- predict(model, x)
# (same as:)
pred <- fitted(model)

# Check accuracy:
table(pred, y)

# compute decision values and probabilities:
pred <- predict(model, x, decision.values = TRUE)
attr(pred, "decision.values")[1:4,]

# visualize (classes by color, SV by crosses):
plot(cmdscale(dist(iris[, -5])),
     col = as.integer(iris[, 5]),
     pch = c("o", "+")[1:150])
```


Description

This generic function tunes hyperparameters of statistical methods using a grid search over supplied parameter ranges.

Usage

```
tune.svm(method, train.x, train.y = NULL, data = list(), validation.x =
  NULL, validation.y = NULL, ranges = NULL, predict.func = predict,
  tunecontrol = tune.control(), ...)
best.tune(...)
```

Arguments

method	either the function to be tuned, or a character string naming such a function.
train.x	either a formula or a matrix of predictors.
train.y	the response variable if train.x is a predictor matrix. Ignored if train.x is a formula.
data	data, if a formula interface is used. Ignored, if predictor matrix and response are supplied directly.
validation.x	an optional validation set. Depending on whether a formula interface is used or not, the response can be included in validation.x or separately specified using validation.y.
validation.y	if no formula interface is used, the response of the (optional) validation set.
ranges	a named list of parameter vectors spanning the sampling space. The vectors will usually be created by seq.
predict.func	optional predict function, if the standard predict behavior is inadequate.
tunecontrol	object of class "tune.control", as created by the function tune.control(). If omitted, tune.control() gives the defaults.
...	Further parameters passed to the training functions.

Details

As performance measure, the classification error is used for classification, and the mean squared error for regression. It is possible to specify only one parameter combination (i.e., vectors of length 1) to obtain an error estimation of the specified type (bootstrap, cross-classification, etc.) on the given data set. For convenience, there are several tune.foo() wrappers defined, e.g., for nnet(), randomForest(), rpart(), svm(), and knn().

Cross-validation randomizes the data set before building the splits which—once created—remain constant during the training process. The splits can be recovered through the train.ind component of the returned object.

Value

For `tune`, an object of class `tune`, including the components:

<code>best.parameters</code>	a 1 x k data frame, k number of parameters.
<code>best.performance</code>	best achieved performance.
<code>performances</code>	if requested, a data frame of all parameter combinations along with the corresponding performance results.
<code>train.ind</code>	list of index vectors used for splits into training and validation sets.
<code>best.model</code>	if requested, the model trained on the complete training data using the best parameter combination.

`best.tune()` returns the best model detected by `tune`.

Author(s)

David Meyer
<David.Meyer@R-project.org>

See Also

[tune.control](#), [plot.tune](#), [tune.svm](#), [tune.wrapper](#)

Examples

```
data(iris)
## tune 'svm' for classification with RBF-kernel (default in svm),
## using one split for training/validation set

obj <- tune(svm, Species~., data = iris,
           ranges = list(gamma = 2^(-1:1), cost = 2^(2:4)),
           tunecontrol = tune.control(sampling = "fix")
           )

## alternatively:
## obj <- tune.svm(Species~., data = iris, gamma = 2^(-1:1), cost = 2^(2:4))

summary(obj)
plot(obj)

## tune 'knn' using a convenience function; this time with the
## conventional interface and bootstrap sampling:
x <- iris[, -5]
y <- iris[, 5]
obj2 <- tune.knn(x, y, k = 1:5, tunecontrol = tune.control(sampling = "boot"))
summary(obj2)
plot(obj2)

## tune 'rpart' for regression, using 10-fold cross validation (default)
data(mtcars)
obj3 <- tune.rpart(mpg~., data = mtcars, minsplit = c(5,10,15))
summary(obj3)
plot(obj3)
```

```
## simple error estimation for lm using 10-fold cross validation
tune(lm, mpg~., data = mtcars)
```

tune.control	<i>Control Parameters for the Tune Function</i>
--------------	---

Description

Creates an object of class `tune.control` to be used with the `tune` function, containing various control parameters.

Usage

```
tune.control(random = FALSE,
             nrepeat = 1,
             repeat.aggregate = mean,
             sampling = c("cross", "fix", "bootstrap"),
             sampling.aggregate = mean,
             sampling.dispersion = sd,
             cross = 10,
             fix = 2/3,
             nboot = 10,
             boot.size = 9/10,
             best.model = TRUE,
             performances = TRUE,
             rough.cross = 0,
             error.fun = NULL)
```

Arguments

<code>random</code>	if an integer value is specified, random parameter vectors are drawn from the parameter space.
<code>nrepeat</code>	specifies how often training shall be repeated.
<code>repeat.aggregate</code>	function for aggregating the repeated training results.
<code>sampling</code>	sampling scheme. If <code>sampling = "cross"</code> , a cross-times cross validation is performed. If <code>sampling = "boot"</code> , <code>nboot</code> training sets of size <code>boot.size</code> (part) are sampled (with replacement) from the supplied data. If <code>sampling = "fix"</code> , a single split into training/validation set is used, the training set containing a <code>fix</code> part of the supplied data. Note that a separate validation set can be supplied via <code>validation.x</code> and <code>validation.y</code> . It is only used for <code>sampling = "boot"</code> and <code>sampling = "fix"</code> ; in the latter case, <code>fix</code> is set to 1.
<code>sampling.aggregate, sampling.dispersion</code>	functions for aggregating the training results on the generated training samples (default: mean and standard deviation).
<code>cross</code>	number of partitions for cross-validation.
<code>fix</code>	part of the data used for training in fixed sampling.
<code>nboot</code>	number of bootstrap replications.

<code>boot.size</code>	size of the bootstrap samples.
<code>best.model</code>	if TRUE, the best model is trained and returned (the best parameter set is used for training on the complete training set).
<code>performances</code>	if TRUE, the performance results for all parameter combinations are returned.
<code>rough.cross</code>	integer.
<code>error.fun</code>	function returning the error measure to be minimized. It takes two arguments: a vector of true values and a vector of predicted values. If NULL, the misclassification error is used for categorical predictions and the mean squared error for numeric predictions.

Value

An object of class "tune.control" containing all the above parameters (either the defaults or the user specified values).

Author(s)

David Meyer
<David.Meyer@R-project.org>

See Also

[tune](#)

Index

*Topic **classification**

predict.gtsvm, 4
svm, 5

*Topic **functions**

load.svmlight, 1

*Topic **models**

plot.tune, 3
tune, 9
tune.control, 11

*Topic **plot**

plot.gtsvm, 2

*Topic **svm**

plot.gtsvm, 2

best.tune (tune), 9

<http://svmlight.joachims.org/>, 2

<http://ttic.uchicago.edu/~textasciitildecotter/projects/gtsvm/>,
5, 8

load.svmlight, 1

Matrix, 5, 6

matrix.csr, 5, 6, 8

plot.gtsvm, 2

plot.tune, 3, 10

predict.gtsvm, 4, 8

print.gtsvm (svm), 5

print.summary.gtsvm (svm), 5

print.summary.tune (tune), 9

print.tune (tune), 9

simple_triplet_matrix, 5, 6

summary.gtsvm (svm), 5

summary.tune (tune), 9

svm, 3, 5, 5

tune, 4, 9, 12

tune.control, 10, 11

tune.svm, 10

tune.wrapper, 10